

# **Machine Learning in Mushroomia**

Computer Science M148

Takao Oba

<https://github.com/toba717/DataScience/tree/main/Project3>

## **I. Introduction**

Mushroomia is a land filled with diverse fauna and flora, but it's particularly well-known for its vast variety of mushrooms. Unfortunately, not all of the mushrooms in Mushroomia are edible, and it can be challenging to differentiate between the ones that are and the ones that aren't. To solve this problem, I have employed Shroomster Pro Max, an advanced data collection device that allows us to gather various data points on any mushroom found in the wild, and The National Archives on Mushrooms, a dataset collected by the Mushroomia government over the years.

Our goal is to construct and train machine learning models on The National Archives on Mushrooms dataset to determine which mushrooms I encounter on our journey are poisonous. I have done so through a systematic process of data loading and exploration, data pre-processing, data augmentation, statistical hypothesis testing, constructing models through non-ensemble and ensemble methods, and tuning the hyperparameters.

## **II. Methodology**

### **A. Data Loading, Splitting, Exploration, and Visualization**

The training dataset and testing dataset was loaded into Google Colab from Google Drive. First, the dataset was split into features and labels, where the “class” column represented the label or the variable in which the machine learning model aims to predict. The rest of the columns were used as tentative predictor variables. The testing data was also split into features and labels as well. Furthermore, the label or the “class” column was transformed from categorical variables ('p' and 'e') to numeric variables (1 and 0) using the map function. Additionally, I ensured that there was no skewness to the label values as this can potentially introduce bias to the model. Next, when scoping each of the predictors, the majority of them were categorical and only three of them were numerical. I visualized the numerical predictors through bar graphs and categorical variables as pie charts. Overall, I ensured that the predictors and the labels did not have any major flaws by first visually inspecting the shape and distribution.

### **B. Data Pre-Processing**

The data must be processed prior to analysis. First, the missing values were identified by checking the number of missing values in the dataset for training and testing separately. The 'veil-type' predictor was identified to have 100% missing values in the testing dataset, and thus it was removed from both the training and testing datasets as it did not provide valuable information.

Overall, to process the data into the best way possible for the model to be used, I constructed a

transformation pipeline. First, to handle the missing values or NaN values, an imputation strategy was used where missing values were replaced by the mode of the corresponding column. This strategy was chosen as it is a simple and effective way to handle missing values in categorical data.

Next, a numeric pipeline was built using the StandardScaler method to scale the numerical features in the dataset. This way, the numeric value's difference in magnitude will not create bias within the model. Finally, the OneHotEncoder method was used to convert categorical data into numerical data suitable for machine learning algorithms. Many of the predictor variables consisted of more than two distinct values. It will be dangerous to assign numeric values with difference in magnitude to these characters as these numbers will possess meanings. The prepared data was then transformed using the built pipeline. Similarly, the testing data was run through this pipeline as well. Overall, I pre-processed my data this way so that the machine learning model can better capture the fed data and further will not run into errors from missing values.

## **C. Data Augmentation**

I aimed to aid the machine learning model through data augmentation or increasing the size and dimension of the dataset. The first new feature we are creating is called "stem\_vol". This feature is constructed by multiplying the "stem-height" and "stem-width" columns of the dataset. This new feature represents the volume of the stem of the plant. This can be useful as certain mushrooms can have a very long stem, but can be skinny and vice versa. Overall, we would like to consider the overall volume so that we are not leaning towards one aspect of the dimension.

The second new feature we are creating is called "diam\_stem\_ratio". We obtain this new feature through dividing the "cap-diameter" column by the "stem-width" column. This new feature represents the ratio of the diameter of the cap to the overall width of the stem. This predictor illustrates the overall shape of the mushroom while simultaneously considering the cap as well as the stem. All in all, I expect that these two features that involve distinct viewpoints of the mushroom can present invaluable insights regarding the class of the mushroom.

## **D. Statistical Hypothesis Testing**

I utilized a subset of the original dataset to perform statistical hypothesis testing. The chosen subset was the numerical predictors: "cap-diameter", "stem-height", and "stem-width". The results of the logistic regression model showed that all three numerical variables had a statistically significant relationship with the class variable. Two of the variables had p-values less than 0.05 and "stem-height" had a p-value that was very close to 0.05. I should therefore include all of these variables when constructing a model.

Furthermore, I noticed that the pseudo R-squared value is 0.02134 which shows that the model

was able to explain a very small amount of variance. This indicates that there may be other factors that contribute to the classification that assess if a mushroom is poisonous or not that is not captured by these three variables. Other categorical variables that I have excluded when performing the hypothesis testing should be included when constructing the machine learning models. This is intuitive as there are definitely more components that we must consider to assess the class of a mushroom than its physical dimension.

## **E. Models of your choice**

I considered multiple models including logistic regression, decision tree, and naive bayes. Logistic Regression was initially used to transform the linear combination of the predictor variables into probability estimates of the binary outcome. Next Decision Trees were considered because they can handle both categorical and continuous data and are simple to understand and interpret. The pruning technique was used to reduce the depth of the tree and make the model less complex. However, pruning the tree led to a drop in accuracy score, highlighting the trade-off between computational cost and performance.

The last non-ensemble method I considered was Naive Bayes because it assumes that the features are independent of each other and that each feature contributes equally to the probability of the class. The Gaussian Naive Bayes model did not yield a high accuracy rate because the normality assumption likely does not hold in these given features. The Bernoulli Naive Bayes model provided a higher accuracy rate because it is most appropriate when working with binary or boolean data, and the data frame had more features involving binary values. Overall, the two distinct Naive Bayes models did not perform very well due to the data having both binary and continuous values, and the overall model assumption that the predictors are independent could have been violated. From the three models that were considered, the decision tree without pruning has performed the best.

## **F. Ensemble Method**

The ensemble method I chose was Random Forest Classifier. The model involves constructing a large number of decision trees on different subsets of the input data and then aggregating their predictions to come up with a prediction. It is an ensemble method because it combines multiple decision trees to generate a more robust and less sensitive model. Random forest showed that it is a great fit for this dataset as it can handle a large number of features and can detect complex relationships between them. Furthermore, by the nature of random forest models, it reduces the risk of overfitting which inevitably implies an increase in testing accuracy. Overall, the random forest classification model has outperformed all of the models that we have constructed thus far.

## **G. Hyper-parameter Tuning**

I aimed to identify the best model through hyper-parameter tuning and finding the best set of hyperparameters. The grid search method is used to explore all possible combinations of hyperparameters by defining a grid of parameter values to search over.

For each model, the corresponding hyperparameters were specified as a list of values to explore using the grid search method. The models were then fitted on the training data, and the function chose the best hyperparameters based on the cross-validation results. The best model was then determined based on the best accuracy score achieved by the model during the cross-validation.

The best hyperparameters found through the evaluations are the following:

- ☐ **Decision Tree Classifier:** max\_depth=30, splitter='random'
- ☐ **Naive Bayes Classifier:** alpha=0.5, force\_alpha=True
- ☐ **Random Forest Classifier:** max\_depth=30, n\_estimators=10

### III. Results

Implementing the best hyper-parameters that was achieved previously, the following was achieved:

- ☐ **Decision Tree Classifier:**
  - ☐ Accuracy: 0.432019
  - ☐ Precision: 0.792294
  - ☐ Recall Score: 0.414434
  - ☐ F1 Score: 0.544205
- ☐ **Naive Bayes Classifier:**
  - ☐ Accuracy: 0.498066
  - ☐ Precision: 1
  - ☐ Recall Score: 0.460228
  - ☐ F1 Score: 0.63051
- ☐ **Random Forest Classifier:**
  - ☐ Accuracy: 0.459285
  - ☐ Precision: 0.956522
  - ☐ Recall Score: 0.4394778
  - ☐ F1 Score: 0.602250

Note that above values can fluctuate for random forest classifier based on the specific seed value that is used.

The Naive Bayes Classifier has yielded the best result at classifying whether a mushroom is poisonous or not. It had yielded the highest value for all four of the evaluation metrics. Overall, the three models tend to have a higher precision score compared to the accuracy score, meaning that the model is correctly classifying positive instances more accurately compared to the negative cases.

These values were obtained through utilizing cross validation with K fold splits of 10. This strategy was chosen as it can provide a good trade off in terms of bias and variance. Furthermore, this method is computationally efficient compared to other strategies.

## **IV. Conclusion**

I look forward to using the Naive Bayes Classifier (with hyperparameters  $\alpha = 0.5$  and `for_alpha = True`) during my adventure through Mushroomia. This is because the overall accuracy metrics was the highest amongst the model indicating that it can capture the true class of the mushroom more accurately so that I can identify which mushroom is poisonous and which one is not.

However, we must consider the limitations of this project. The dataset from the National Archives on Mushrooms which may not necessarily capture the entire picture of the current notion as it may be out of date or in reality may involve more predictors than given in the dataset. For example, if there are more predators or mushroom eaters in a specific region, there is a high chance that the mushrooms are not poisonous. Moreover, it is difficult to truly identify if a mushroom is poisonous or not and thus when recording and collecting the data, there could have been misclassification which can heavily impact the machine learning model. Furthermore, the results of our visualization were limited by the quality of the available data. There were a limited number of numerical and categorical variables to accurately visualize the data.

Regarding next steps, numerical variables can be assessed amongst each other through considering multicollinearity and further considering association for categorical variables. Additionally, locations can be addressed as well so that the machine learning model can be used in other contexts of mushroom classification as well.

You are exploring the wilderness of *Mushroomia*, a land populated by a plethora of diverse fauna and flora. In particular, *Mushroomia* is known for its unparalleled variety in mushrooms. However, not all the mushrooms in *Mushroomia* are edible. As you make your way through *Mushroomia*, you would like to know which mushrooms are edible, in order to forage for supplies for your daily mushroom soup.

You have access to:

- *Shroomster Pro Max™* - a state of the art data collection device, developed by *Mushroomia*, that allows you to collect various data points about any mushroom you encounter in the wild
- *The National Archives on Mushrooms* - a dataset collected over the years by the government of *Mushroomia*

To address this problem, you decide to use the skills you learnt in CSM148 and train machine learning models on the *The National Archives on Mushrooms* in order to use your *Shroomster Pro Max™* to determine whether the mushrooms you encounter on your adventure can be added to your daily mushroom soup.

This project will be more unstructured than the previous two projects in order to allow you to experience how data science problems are solved in practice. There are two parts to this project: a Jupyter Notebook with your code (where you explore, visualize, process your data and train machine learning models) and a report (where you explain the various choices you make in your implementation and analyze the final performance of your models).

## 1. Loading and Viewing Data

```
[10] # mount to google drive and retrieve files
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

[11] # importing neccessary modules
import numpy as np
import pandas as pd

%matplotlib inline
import matplotlib.pyplot as plt

[12] # download the necessary files that is stored in the google drive
train = pd.read_csv("/content/drive/MyDrive/CS148_data/project3/mushroom_train.csv", sep = ";")
test = pd.read_csv("/content/drive/MyDrive/CS148_data/project3/mushroom_test.csv", sep = ";")
```

13] # viewing the training data

train.head()

	class	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	...	stem-root	stem-surface	stem-color	veil-type	veil-color	has-ring	ring-type	spore-print-color	habitat	season	
0	p	15.26	x	g	o		f	e	NaN	w	16.95	...	s	y	w	u	w	t	g	NaN	d	w
1	p	16.60	x	g	o		f	e	NaN	w	17.99	...	s	y	w	u	w	t	g	NaN	d	u
2	p	14.07	x	g	o		f	e	NaN	w	17.80	...	s	y	w	u	w	t	g	NaN	d	w
3	p	14.17	f	h	e		f	e	NaN	w	15.77	...	s	y	w	u	w	t	p	NaN	d	w
4	p	14.64	x	h	o		f	e	NaN	w	16.53	...	s	y	w	u	w	t	p	NaN	d	w

5 rows x 21 columns

```

[14] # viewing the testing data
test.head()

```

	class	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	...	stem-root	stem-surface	stem-color	veil-type	veil-color	has-ring	ring-type	spore-print-color	habitat	season	
0	p	2.50	b	NaN	k		f	a	NaN	k	8.42	...	NaN	NaN	g	NaN	w	f	f	k	g	u
1	p	3.07	b	NaN	k		f	a	NaN	n	7.24	...	NaN	NaN	n	NaN	w	f	f	k	g	a
2	p	3.30	b	NaN	n		f	a	NaN	n	10.22	...	NaN	NaN	n	NaN	w	f	f	k	g	u
3	p	3.49	b	NaN	k		f	a	NaN	k	11.00	...	NaN	NaN	n	NaN	w	f	f	k	g	a
4	p	2.79	b	NaN	n		f	a	NaN	n	6.97	...	NaN	NaN	g	NaN	w	f	f	k	g	u

5 rows x 21 columns

## 2. Splitting Data into Features and Labels

```
[15] # This will be the label that the model aims to predict
label = train['class']
label
```

```
0      p
1      p
2      p
3      p
4      p
...
50208   p
50209   p
50210   p
50211   p
50212   p
Name: class, Length: 50213, dtype: object
```

```
[16] # This will be the tentative predictor variables
train_dropped = train.drop(['class'], axis = 1)
train_dropped
```

	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	stem-width	stem-root	stem-surface	stem-color	veil-type	veil-color	has-ring	ring-type	spore-print-color	habitat	season	
0	15.26	x	g	o		f	e	NaN	w	16.95	17.09	s	y	w	u	w	t	g	NaN	d	w
1	16.60	x	g	o		f	e	NaN	w	17.99	18.19	s	y	w	u	w	t	g	NaN	d	u
2	14.07	x	g	o		f	e	NaN	w	17.80	17.74	s	y	w	u	w	t	g	NaN	d	w
3	14.17	f	h	e		f	e	NaN	w	15.77	15.98	s	y	w	u	w	t	p	NaN	d	w
4	14.64	x	h	o		f	e	NaN	w	16.53	17.20	s	y	w	u	w	t	p	NaN	d	w
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
50208	1.18	s	s	y		f	f	f	f	3.93	6.22	NaN	NaN	y	NaN	NaN	f	f	NaN	d	a
50209	1.27	f	s	y		f	f	f	f	3.18	5.43	NaN	NaN	y	NaN	NaN	f	f	NaN	d	a
50210	1.27	s	s	y		f	f	f	f	3.86	6.37	NaN	NaN	y	NaN	NaN	f	f	NaN	d	u
50211	1.24	f	s	y		f	f	f	f	3.56	5.44	NaN	NaN	y	NaN	NaN	f	f	NaN	d	u
50212	1.17	s	s	y		f	f	f	f	3.25	5.45	NaN	NaN	y	NaN	NaN	f	f	NaN	d	u

50213 rows x 20 columns

```
# predictors for the testing dataset
test_dropped = test.drop(['class'], axis = 1)
test_dropped
```

	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	stem-width	stem-root	stem-surface	stem-color	veil-type	veil-color	has-ring	ring-type	spore-print-color	habitat	season	
0	2.50	b	NaN	k		f	a	NaN	k	8.42	2.46	NaN	NaN	g	NaN	w	f	f	k	g	u
1	3.07	b	NaN	k		f	a	NaN	n	7.24	2.41	NaN	NaN	n	NaN	w	f	f	k	g	a
2	3.30	b	NaN	n		f	a	NaN	n	10.22	2.53	NaN	NaN	n	NaN	w	f	f	k	g	u
3	3.49	b	NaN	k		f	a	NaN	k	11.00	2.81	NaN	NaN	n	NaN	w	f	f	k	g	a
4	2.79	b	NaN	n		f	a	NaN	n	6.97	2.37	NaN	NaN	g	NaN	w	f	f	k	g	u
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
10851	52.41	o	y	y		f	p	NaN	y	5.47	25.02	NaN	k	k	NaN	NaN	f	f	NaN	d	u
10852	54.81	o	y	y		f	p	NaN	y	6.67	22.15	NaN	k	k	NaN	NaN	f	f	NaN	d	s
10853	49.95	o	y	y		f	p	NaN	y	6.43	26.35	NaN	k	n	NaN	NaN	f	f	NaN	d	u
10854	53.16	o	y	y		f	p	NaN	y	6.99	40.29	NaN	k	k	NaN	NaN	f	f	NaN	d	s
10855	49.78	o	y	y		f	p	NaN	y	5.77	18.26	NaN	k	k	NaN	NaN	f	f	NaN	d	u

10856 rows x 20 columns

```
[18] # the target variable for the testing data
label_test = test['class']
print(label_test)
```

```
0      p
1      p
2      p
3      p
4      p
...
10851   e
10852   e
10853   e
10854   e
10855   e
Name: class, Length: 10856, dtype: object
```

```
[19] # changing the labels to numeric values
label = label.map({"p":1, "e":0})
label_test = label_test.map({"p":1, "e":0})
```



### ▼ 3. Data Exploration, Visualization, and Data Processing

```
✓ [20] # view the dimensions of the training and testing data
0a
print(train.shape)
print(test.shape)
print(train_dropped.shape)
print(test_dropped.shape)
```

```
(50213, 21)
(10856, 21)
(50213, 20)
(10856, 20)
```

```
✓ [21] # look at the type of each predictors
0a
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 50213 entries, 0 to 50212
Data columns (total 21 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   class                 50213 non-null  object
 1   cap-diameter          50213 non-null  float64
 2   cap-shape             50213 non-null  object
 3   cap-surface          37915 non-null  object
 4   cap-color            50213 non-null  object
 5   does-bruise-or-bleed  50213 non-null  object
 6   gill-attachment      42447 non-null  object
 7   gill-spacing         31064 non-null  object
 8   gill-color           50213 non-null  object
 9   stem-height          50213 non-null  float64
10  stem-width           50213 non-null  float64
11  stem-root            7413 non-null   object
12  stem-surface         19912 non-null  object
13  stem-color           50213 non-null  object
14  veil-type            3177 non-null   object
15  veil-color           6297 non-null   object
16  has-ring             50213 non-null  object
17  ring-type            48448 non-null  object
18  spore-print-color     4532 non-null   object
19  habitat              50213 non-null  object
20  season              50213 non-null  object
dtypes: float64(3), object(18)
memory usage: 8.0+ MB
```

```
✓ [22] # view the amount of NA values in this dataset for training
0a
train.isnull().sum()
```

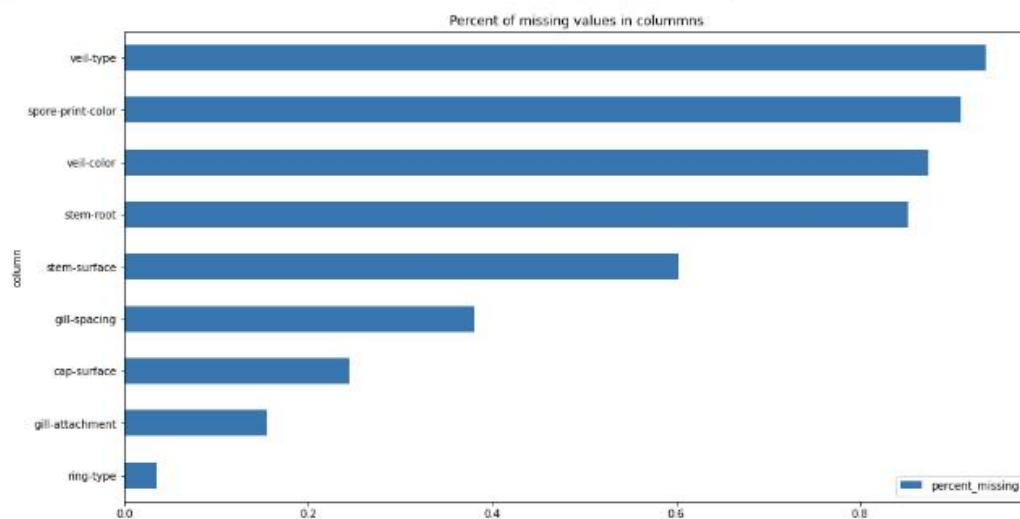
```
class                0
cap-diameter         0
cap-shape            0
cap-surface         12298
cap-color            0
does-bruise-or-bleed 0
gill-attachment      7766
gill-spacing        19149
gill-color           0
stem-height          0
stem-width           0
stem-root           42800
stem-surface        30301
stem-color           0
veil-type           47036
veil-color          43916
has-ring             0
ring-type           1765
spore-print-color    45681
habitat              0
season              0
dtype: int64
```

```
# Visualize the amount of missing values for each column

import pylab

def plot_missing_values(df):
    data = [(col, df[col].isnull().sum() / len(df))
            for col in df.columns if df[col].isnull().sum() > 0]
    col_names = ['column', 'percent_missing']
    missing_df = pd.DataFrame(data, columns=col_names).sort_values('percent_missing')
    pylab.rcParams['figure.figsize'] = (15, 8)
    missing_df.plot(kind='barh', x='column', y='percent_missing');
    plt.title('Percent of missing values in columns');

plot_missing_values(train)
```

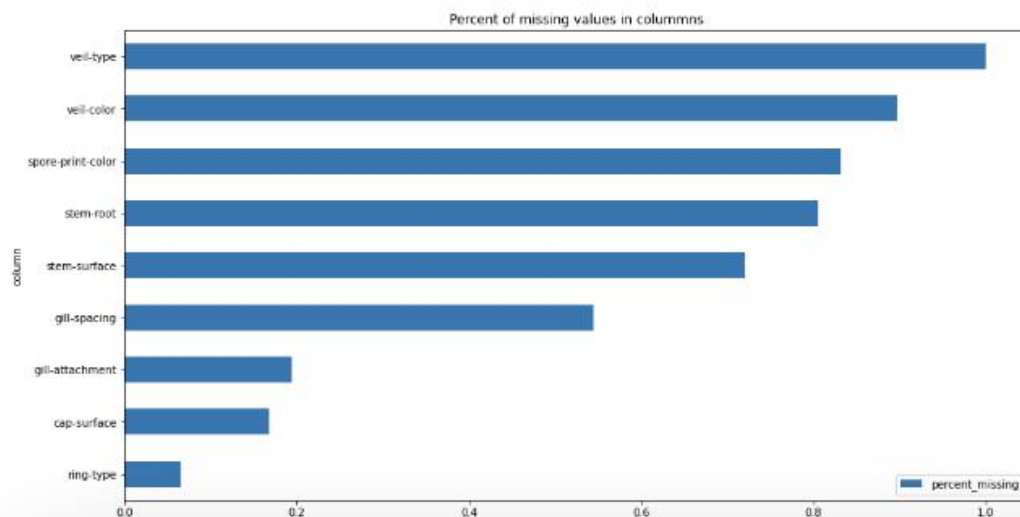


```
# view the amount of missing values for testing
test.isnull().sum()

# columns with missing values are the following
# missing_vals = ["cap-surface", "gill-attachment", "gill-spacing", "stem-root", "stem-surface", "veil-type", "veil-color", "ring-type", "spore-print-color"]
```

```
class          0
cap-diameter    0
cap-shape       0
cap-surface    1822
cap-color       0
does-bruise-or-bleed 0
gill-attachment 2118
gill-spacing    5914
gill-color      0
stem-height     0
stem-width      0
stem-root       8738
stem-surface    7823
stem-color      0
veil-type      10856
veil-color      9740
has-ring        0
ring-type       706
spore-print-color 9034
habitat         0
season          0
dtype: int64
```

```
[25] # Visualize the amount of missing values for each column
plot_missing_values(test)
```




Note that 100% of the values in the "veil-type" column is missing (there are 10856 observations in the testing data set and 10856 of them are missing values for this specific column). Thus, we have that this column does not provide us with valuable information, and can be removed from both the training and testing dataset.

```
[26] # drop "veil-type"
train_dropped = train_dropped.drop(["veil-type"], axis = 1)
test_dropped = test_dropped.drop(["veil-type"], axis = 1)
```


```
[27] # find the specifics for each of the numerical predictors
train_dropped.describe()
```

	cap-diameter	stem-height	stem-width
count	50213.000000	50213.000000	50213.000000
mean	6.245186	6.600561	10.763191
std	4.542552	3.221714	7.744992
min	0.380000	1.200000	0.520000
25%	3.290000	4.680000	4.720000
50%	5.540000	5.900000	9.130000
75%	8.100000	7.600000	15.210000
max	58.890000	33.920000	58.950000

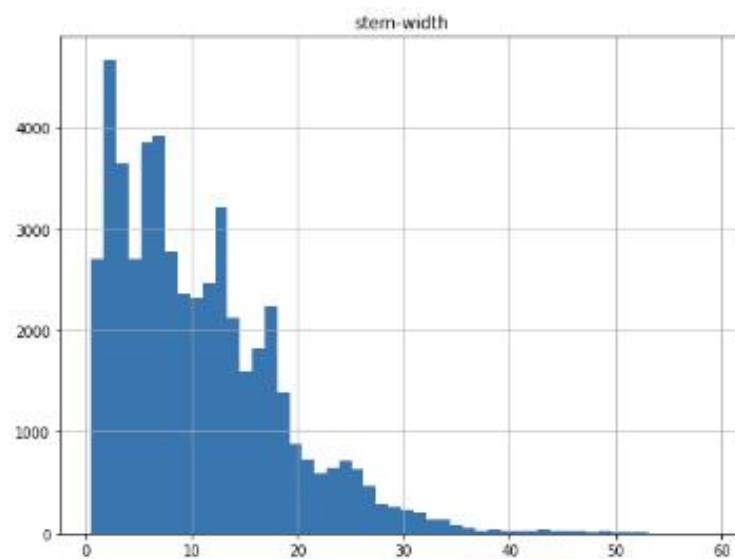
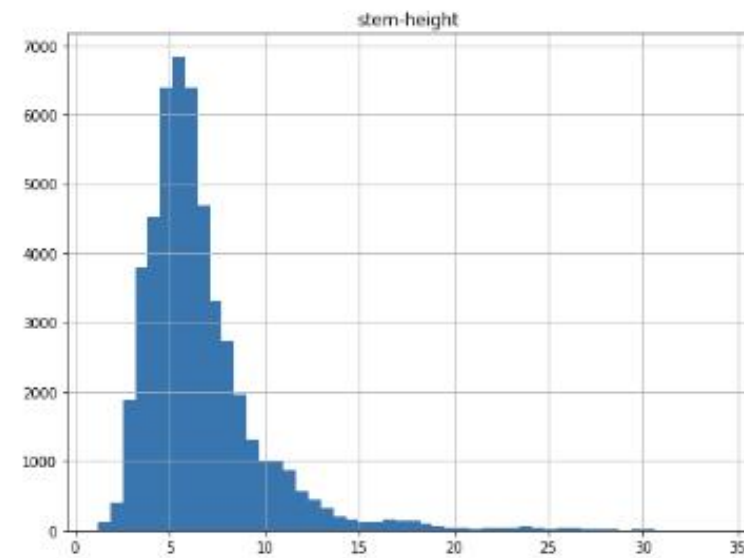
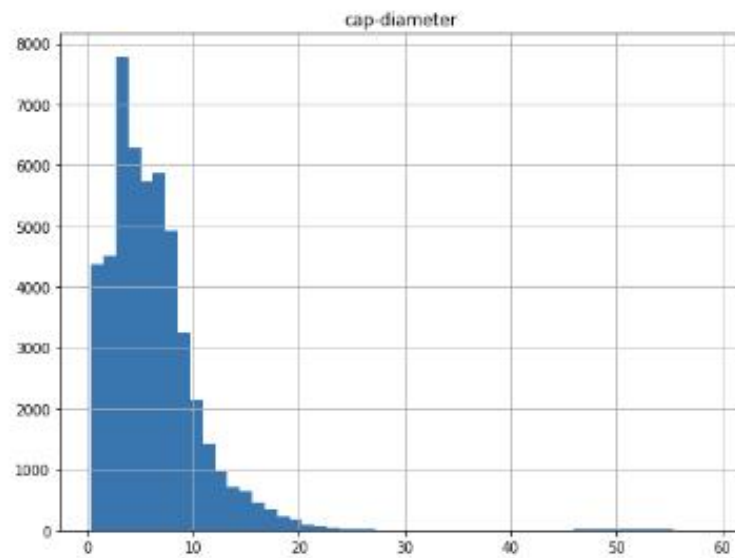
```
[28] # find correlation of the numeric predictors
train_dropped.corr()
import seaborn as sns
sns.heatmap(train_dropped.corr())
plt.show()
```



```
# visualization for the target variable  
label.hist(bins = 50, figsize = (20,15))  
plt.show()
```




```
# visualization for numeric, predictor variables
train_dropped.hist(bins = 50, figsize = (20,15))
plt.show()
```



From above, we can see that there is a great amount of samples that are both poisonous and not poisonous. We can conclude that our classifier will have a sufficiently balanced dataset to adequately classify the two (i.e. there is no high skew in the two values as this can be problematic when considering accuracy rates).


```
# visualization of the categorical variables
temp = (train_dropped.drop(["cap-diameter", "stem-height", "stem-width"], axis = 1)).apply(pd.value_counts).fillna(0)
temp1 = temp.iloc[:,0:6]
fig = temp1.plot(kind = 'pie',subplots = True, figsize = (30,30), legend = None)

plt.show()
```




```
[32] # next 5 predictors
temp2 = temp.iloc[:, 6:11]
fig = temp2.plot(kind = 'pie',subplots = True, figsize = (30,30), legend = None)

plt.show()
```



```
[33] # next 5 predictors
temp3 = temp.iloc[:,11:]
fig = temp3.plot(kind = 'pie',subplots = True, figsize = (30,30), legend = None)

plt.show()
```





```
✓ [34] # below shows whether we should scale or encode any of the variables
```

```
09 for col in train_dropped:
    print(col)
    print(train_dropped[col].unique())

cap-diameter
[15.26 16.6 14.07 ... 21.8 20.42 22.71]
cap-shape
['x' 'f' 'p' 'b' 'c' 's' 'o']
cap-surface
['g' 'h' nan 't' 'y' 'e' 's' 'l' 'd' 'w' 'i' 'k']
cap-color
['o' 'e' 'n' 'g' 'r' 'w' 'y' 'p' 'u' 'b' 'l' 'k']
does-bruise-or-bleed
['f' 't']
gill-attachment
['e' nan 'a' 'd' 's' 'x' 'p' 'f']
gill-spacing
[nan 'c' 'd' 'f']
gill-color
['w' 'n' 'p' 'u' 'b' 'g' 'y' 'r' 'e' 'o' 'k' 'f']
stem-height
[16.95 17.99 17.8 ... 17.7 14.62 15.15]
stem-width
[17.09 18.19 17.74 ... 21.78 21.56 22.53]
stem-root
['s' nan 'b' 'r']
stem-surface
['y' nan 's' 'k' 'i' 'h' 't' 'g']
stem-color
['w' 'y' 'n' 'u' 'b' 'l' 'r' 'p' 'e' 'k' 'g' 'o']
veil-color
['w' 'y' nan 'n' 'e' 'u']
has-ring
['t' 'f']
ring-type
['g' 'p' 'e' 'l' 'f' 'm' nan 'r' 'z']
spore-print-color
[nan 'w' 'p' 'k' 'r' 'u' 'n']
habitat
['d' 'm' 'g' 'h' 'l' 'p' 'w' 'u']
season
['w' 'u' 'a' 's']
```

```
✓ [35] # since "stem-width" is in mm and "stem-height", "cap-diameter" are in cm, we aim to convert "stem-height" and "cap-diameter"
09 # if we aim to convert "stem-width", since it is a small value, we may lose precision
```

```
# for training data
train_dropped['cap-diameter'] = train_dropped['cap-diameter']*100
train_dropped['stem-height'] = train_dropped['stem-height']*100

# for testing data
test_dropped['cap-diameter'] = test_dropped['cap-diameter']*100
test_dropped['stem-height'] = test_dropped['stem-height']*100
```

```
✓ [36] # build numeric pipeline
09 from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer, make_column_transformer

train_num = train_dropped[["cap-diameter", "stem-height", "stem-width"]]

num_pipeline = Pipeline([
    ('std_scaler', StandardScaler())
])

numerical_features = list(train_num)
```

```
✓ [37] # generate variable that will capture features that are used for one hot encoding
09 OHE = list(train_dropped.drop(["cap-diameter", "stem-height", "stem-width"], axis = 1))
```

✓  
01

```
[38] # for training data
# use to impute missing values with mode

# first store column names in variable
column_names = train_dropped.columns

# simple imputation using mode
imp = SimpleImputer(missing_values = np.nan, strategy = 'most_frequent')

# fit the simple imputation strategy
imputed_train_dropped = imp.fit_transform(train_dropped)

# revert back to data frame
imputed_train_dropped = pd.DataFrame(imputed_train_dropped, columns = column_names).reset_index(drop = 'index')
imputed_train_dropped
```

	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	stem-width	stem-root	stem-surface	stem-color	veil-color	has-ring	ring-type	spore-print-color	habitat	season
0	1526.0	x	g	o	f	e	c	w	1695.0	17.09	s	y	w	w	t	g	k	d	w
1	1660.0	x	g	o	f	e	c	w	1799.0	18.19	s	y	w	w	t	g	k	d	u
2	1407.0	x	g	o	f	e	c	w	1780.0	17.74	s	y	w	w	t	g	k	d	w
3	1417.0	f	h	e	f	e	c	w	1577.0	15.98	s	y	w	w	t	p	k	d	w
4	1464.0	x	h	o	f	e	c	w	1653.0	17.2	s	y	w	w	t	p	k	d	w
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
50208	118.0	s	s	y	f	f	f	f	393.0	6.22	b	s	y	w	f	f	k	d	a
50209	127.0	f	s	y	f	f	f	f	318.0	5.43	b	s	y	w	f	f	k	d	a
50210	127.0	s	s	y	f	f	f	f	386.0	6.37	b	s	y	w	f	f	k	d	u
50211	124.0	f	s	y	f	f	f	f	356.0	5.44	b	s	y	w	f	f	k	d	u
50212	117.0	s	s	y	f	f	f	f	325.0	5.45	b	s	y	w	f	f	k	d	u

50213 rows x 19 columns

✓  
01

```
# for testing data
# use to impute missing values with mode

# first store column names in variable
column_names = test_dropped.columns

# simple imputation using mode
imp = SimpleImputer(missing_values = np.nan, strategy = 'most_frequent')

# fit the simple imputation strategy
imputed_test_dropped = imp.fit_transform(test_dropped)

# revert back to data frame
imputed_test_dropped = pd.DataFrame(imputed_test_dropped, columns = column_names).reset_index(drop = 'index')
imputed_test_dropped
```



	cap-diameter	cap-shape	cap-surface	cap-color	does-bruise-or-bleed	gill-attachment	gill-spacing	gill-color	stem-height	stem-width	stem-root	stem-surface	stem-color	veil-color	has-ring	ring-type	spore-print-color	habitat	season
0	250.0	b	t	k	f	a	c	k	842.0	2.46	f	f	g	w	f	f	k	g	u
1	307.0	b	t	k	f	a	c	n	724.0	2.41	f	f	n	w	f	f	k	g	a
2	330.0	b	t	n	f	a	c	n	1022.0	2.53	f	f	n	w	f	f	k	g	u
3	349.0	b	t	k	f	a	c	k	1100.0	2.81	f	f	n	w	f	f	k	g	a
4	279.0	b	t	n	f	a	c	n	697.0	2.37	f	f	g	w	f	f	k	g	u
...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...	...
10851	5241.0	o	y	y	f	p	c	y	547.0	25.02	f	k	k	w	f	f	k	d	u
10852	5481.0	o	y	y	f	p	c	y	667.0	22.15	f	k	k	w	f	f	k	d	s
10853	4995.0	o	y	y	f	p	c	y	643.0	26.35	f	k	n	w	f	f	k	d	u
10854	5316.0	o	y	y	f	p	c	y	699.0	40.29	f	k	k	w	f	f	k	d	s
10855	4978.0	o	y	y	f	p	c	y	577.0	18.26	f	k	k	w	f	f	k	d	u

10856 rows x 19 columns



```
[40] # build fill pipeline
full_pipeline = ColumnTransformer([
    ('num', num_pipeline, numerical_features),
    ('cat', OneHotEncoder(categories = "auto", handle_unknown = "ignore", sparse_output = False), OHE),
])
```

```
[41] # transform training data
train_prepared = full_pipeline.fit_transform(imputed_train_dropped)
print(train_prepared.shape)

(50213, 112)
```

```
[42] # transform testing data
test_prepared = full_pipeline.fit_transform(imputed_test_dropped)
print(test_prepared.shape)

(10856, 87)
```

```
# as fit_transform outputs an np.array, we aim to convert to a df

# for training data

# we aim to retrieve column names
ohe_encoder = OneHotEncoder(categories = "auto", handle_unknown = "ignore", sparse_output = False)
X_object = (imputed_train_dropped.drop(["cap-diameter", "stem-height", "stem-width"], axis = 1))
ohe_encoder.fit(X_object)

# features that will be introduced through One Hot Encoding
feature_names_ohe = ohe_encoder.get_feature_names_out(OHE)
feature_names_ohe

# convert the np.array to df (this df will not have column names)
train_prepared_df = pd.DataFrame(train_prepared)

# combine numerical features and features that were OHE
features_df = numerical_features + list(feature_names_ohe)
features_df

# rename df
train_prepared_df.set_axis(features_df, axis = 1, inplace = True)

train_prepared_df.head()
```

	cap-diameter	stem-height	stem-width	cap-shape_b	cap-shape_c	cap-shape_f	cap-shape_o	cap-shape_p	cap-shape_s	cap-shape_x	...	habitat_h	habitat_l	habitat_m	habitat_p	habitat_u	habitat_w	season_a	season_s	season_u	season_w
0	1.984546	3.212433	0.816899	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
1	2.279537	3.535246	0.958927	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
2	1.722576	3.476271	0.900825	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
3	1.744590	2.846165	0.673579	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
4	1.848058	3.082067	0.831101	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0

5 rows x 112 columns

```
[44] # for testing data

# we aim to retrieve column names
ohe_encoder = OneHotEncoder(categories = "auto", handle_unknown = "ignore", sparse_output = False)
X_object = (imputed_test_dropped.drop(["cap-diameter", "stem-height", "stem-width"], axis = 1))
ohe_encoder.fit(X_object)

# features that will be introduced through One Hot Encoding
feature_names_ohe = ohe_encoder.get_feature_names_out(OHE)
feature_names_ohe

# convert the np.array to df (this df will not have column names)
test_prepared_df = pd.DataFrame(test_prepared)

# combine numerical features and features that were OHE
features_df = numerical_features + list(feature_names_ohe)
features_df

# rename df
test_prepared_df.set_axis(features_df, axis = 1, inplace = True)

test_prepared_df.head()
```

	cap-diameter	stem-height	stem-width	cap-shape_b	cap-shape_f	cap-shape_o	cap-shape_p	cap-shape_s	cap-shape_x	cap-surface_d	...	spore-print-color_p	spore-print-color_w	habitat_d	habitat_g	habitat_h	habitat_l	season_a	season_s	season_u	season_w
0	-0.881539	0.483594	-1.041207	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
1	-0.804165	0.187381	-1.044440	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0
2	-0.772944	0.935446	-1.036680	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0
3	-0.747152	1.131248	-1.018573	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0
4	-0.842173	0.119603	-1.047027	1.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	1.0	0.0

5 rows x 87 columns

## 5. Data Augmentation (Creating at least 2 New Features)

### "stem\_vol" Predictor

```
[45] # use volume of stem as a new predictor
stem_vol_train = train_prepared_df["stem-height"]*train_prepared_df["stem-width"] # training data
stem_vol_train
```

```
0      2.624232
1      3.390044
2      3.131510
3      1.917116
4      2.561510
...
50208   0.486255
50209   0.731114
50210   0.482526
50211   0.648674
50212   0.713467
Length: 50213, dtype: float64
```

```
[46] # adding feature to the prepared training dataframe
train_prepared_df["stem_vol"] = stem_vol_train
```

```
[47] # do the same for testing data
stem_vol_test = test_prepared_df["stem-height"]*test_prepared_df["stem-width"]
stem_vol_test
```

```
0      -0.503522
1      -0.195708
2      -0.969758
3      -1.152259
4      -0.125227
...
10851  -0.107316
10852   0.010280
10853  -0.008035
10854   0.175111
10855   0.003538
Length: 10856, dtype: float64
```

```
[48] # adding feature to the prepared testing dataframe
test_prepared_df["stem_vol"] = stem_vol_test
```

### "diam\_stem\_ratio" Predictor

```
▶ # ratio between cap-diameter and stem-width

diam_stem_ratio_train = train_prepared_df["cap-diameter"]/train_prepared_df["stem-width"] # training data
diam_stem_ratio_train
```

```
0      2.429367
1      2.377175
2      1.912222
3      2.590032
4      2.223625
...
50208   1.900883
50209   1.590535
50210   1.930858
50211   1.603132
50212   1.628612
Length: 50213, dtype: float64
```

```
[50] # adding feature to the prepared training dataframe
train_prepared_df["diam_stem_ratio"] = diam_stem_ratio_train
```

```
[51] # do the same for testing data

diam_stem_ratio_test = test_prepared_df["cap-diameter"]/test_prepared_df["stem-width"]
diam_stem_ratio_test
```

```
0      0.846651
1      0.769948
2      0.745595
3      0.733528
4      0.804348
...
10851  14.110339
10852  26.798266
10853  11.037903
10854   4.266704
10855 -284.261085
Length: 10856, dtype: float64
```

```
[52] # adding feature to the prepared testing dataframe
test_prepared_df["diam_stem_ratio"] = diam_stem_ratio_test
```

## ▼ 6. Logistic Regression & Statistical Hypothesis Testing

### ▼ Logistic Regression

```
[53] # splitting training data to have a validation set
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV
X_train, X_test, Y_train, Y_test = train_test_split(train_prepared_df, label, test_size = 0.2)
```

```
[54] # fit logistic regression model
from sklearn.linear_model import LogisticRegression
log_model = LogisticRegression(random_state = 0, max_iter = 3000)
log_model_fit = log_model.fit(X_train, Y_train)
log_pred = log_model.predict(X_test)
```

```
[55] # accuracy metrics of logistic regression model
from sklearn import metrics

# the accuracy of the model
print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(log_pred, Y_test)))

# the precision of the model
print("%-12s %f" % ('Precision:', metrics.precision_score(log_pred, Y_test)))

# the recall score of the model
print("%-12s %f" % ('Recall Score:', metrics.recall_score(log_pred, Y_test)))

# the F1 score of the model
print("%-12s %f" % ('F1 Score:', metrics.f1_score(log_pred, Y_test)))
```

```
Accuracy:    0.859803
Precision:   0.875297
Recall Score: 0.884892
F1 Score:    0.880068
```

### ▼ Statistical Hypothesis Testing

```
# the subset with 3 or more features
# define the subset to be the first three columns or the numerical features
train_prepared_df.iloc[:,0:3]
```

	cap-diameter	stem-height	stem-width
0	1.984546	3.212433	0.816899
1	2.279537	3.535246	0.958927
2	1.722576	3.476271	0.900825
3	1.744590	2.846165	0.673579
4	1.848058	3.082067	0.831101
...	...	...	...
50208	-1.115064	-0.828934	-0.586603
50209	-1.095251	-1.061732	-0.688605
50210	-1.095251	-0.850662	-0.567236
50211	-1.101855	-0.943781	-0.687314
50212	-1.117265	-1.040004	-0.686023

50213 rows x 3 columns

```
[57] import statsmodels.api as sm

# run to add_constant to add a constant feature to df that will serve as Y intercept
sm_x = sm.add_constant(train_prepared_df.iloc[:,0:3])

shroom_stats = sm.Logit(label, sm_x)

results_stats = shroom_stats.fit()

print(results_stats.summary())
```

```
Optimization terminated successfully.
Current function value: 0.665016
Iterations 5
```

```
Logit Regression Results
=====
Dep. Variable:          class      No. Observations:          50213
Model:                Logit      Df Residuals:           50209
Method:                MLE        Df Model:                3
Date:                  Sun, 05 Mar 2023    Pseudo R-squ.:         0.02134
Time:                  11:11:00    Log-Likelihood:         -33392.
converged:              True        LL-Null:               -34121.
Covariance Type:        nonrobust    LLR p-value:            0.000
=====
                    coef    std err          z      P>|z|      [0.025      0.975]
=====
const              0.3326      0.009     36.200      0.000      0.315      0.351
cap-diameter      -0.4377      0.020    -21.460      0.000     -0.478     -0.398
stem-height        0.0210      0.011      1.918     0.055     -0.000      0.042
stem-width         0.0512      0.016      3.162     0.002      0.019      0.083
=====
```

## 7. Dimensionality Reduction using PCA

```
[58] # PCA: https://scikit-learn.org/stable/modules/generated/sklearn.decomposition.PCA.html
```

```
[59] # import necessary packages
      from sklearn import decomposition
```

```
[60] # notice the shape of the current training set
      X_train.shape

      (40170, 114)
```

```
[61] # we aim to decompose this through pca
      # first, we decompose through singular value decomposition
      # assign a value to number of components that will be utilized
      pca = decomposition.PCA(n_components = 10)

      pca_train = pca.fit_transform(X_train)
```

```
▶ # notice that this aligns with the number of components
    pca_train.shape

    (40170, 10)
```

```
[63] # perform logistic regression on these components
      log_reg = LogisticRegression(solver = 'liblinear')
      log_reg.fit(pca_train, Y_train)
```

```
▼      LogisticRegression
LogisticRegression(solver='liblinear')
```

```
[64] # now that we have our model from these components, we assess accuracy on validation data
      pca_val = pca.transform(X_test)
      predicted = log_reg.predict(pca_val)

      print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(Y_test, predicted)))

      Accuracy:      0.652096
```

```
[65] # tells us the proportion of the total variance in the data that is explained by each PC.
      print(pca.explained_variance_ratio_)

      # represent the amount of variance in the data that is accounted for by each principal component
      print(pca.singular_values_)

      [9.77369313e-01 4.54321975e-03 2.41092619e-03 1.15340704e-03
       9.21429073e-04 8.27259717e-04 8.02206587e-04 7.76805854e-04
       6.50134638e-04 6.16988100e-04]
      [4505.17409913 307.15960544 223.75574069 154.76527682 138.32898318
       131.06996021 129.07000919 127.01016319 116.19404789 113.19327435]
```

The accuracy has gone down when utilizing principal component analysis.

```
[66] # Now we do the same for the testing data
      pca = decomposition.PCA(n_components = 10)

      pca_test = pca.fit_transform(test_prepared_df)

      pca_test.shape

      (10856, 10)
```



## ▼ 8. Experiment with any 2 other models (Non-Ensemble)

18 [67] # Models: [https://scikit-learn.org/stable/supervised\\_learning.html](https://scikit-learn.org/stable/supervised_learning.html)

### ▼ Model 1: Decision Trees

```
19 # Decision Trees

from sklearn import tree

# call the model
tree_model = tree.DecisionTreeClassifier()

# fit the model with our data
tree_fit = tree_model.fit(X_train, Y_train)

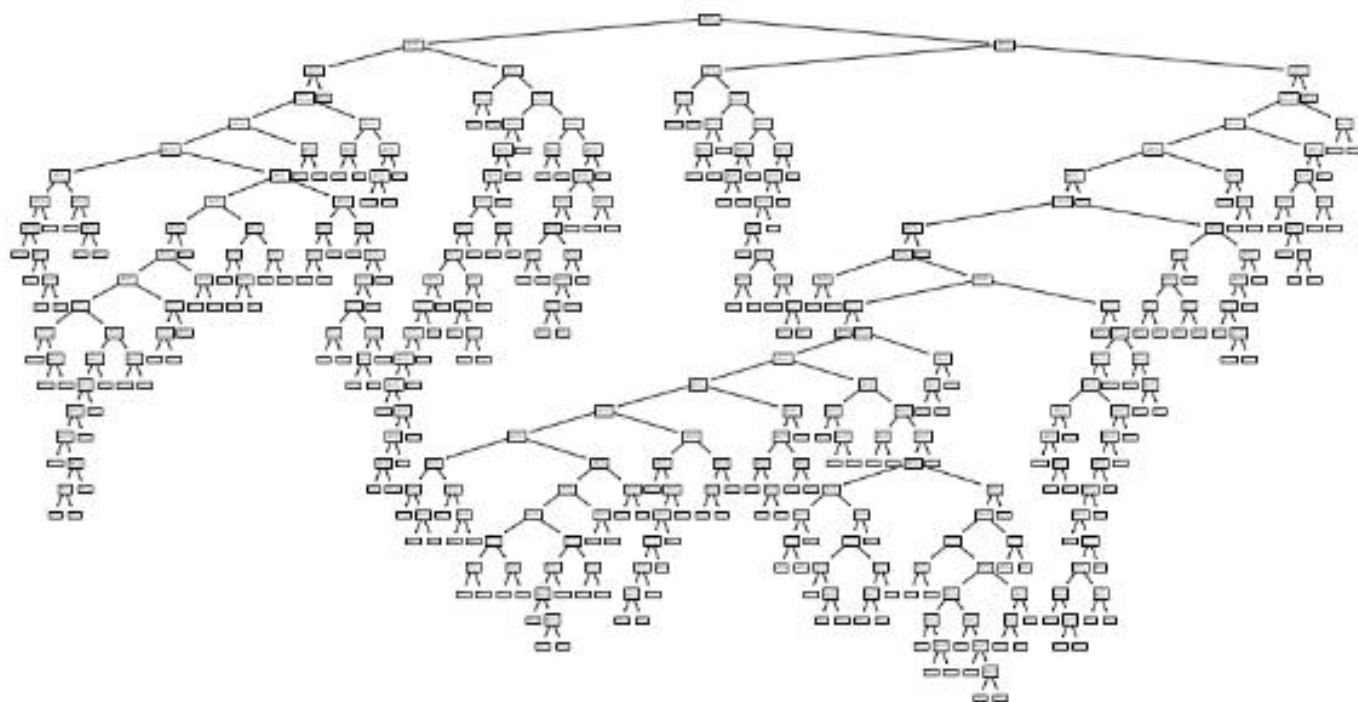
# predict using our model
tree_pred = tree_fit.predict(X_test)

# accuracy score
print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(tree_pred, Y_test)))
```

Accuracy: 0.998208

```
40 [69] # overall image of decision tree
print(tree.plot_tree(tree_model))
```

[Text(0.5163931005141315, 0.9814814814814815, 'x[77] <= 0.5\ngini = 0.487\nsamples = 40170\nvalue = [168.



```

▶ ### pruning decision tree

# find the best max_length to use without negatively affecting the decision tree
max_depth = []
acc_gini = []
acc_entropy = []
for i in range(1,30):
    dtree = tree.DecisionTreeClassifier(criterion='gini', max_depth=i)
    dtree.fit(X_train, Y_train)
    pred = dtree.predict(X_test)
    acc_gini.append(metrics.accuracy_score(Y_test, pred))
    ###
    max_depth.append(i)
d = pd.DataFrame({'acc_gini':pd.Series(acc_gini),
                  'max_depth':pd.Series(max_depth)})

# plot to find the max_depth value when elbow occurs
plt.plot('max_depth','acc_gini', data=d, label='gini')
plt.xlabel('max_depth')
plt.ylabel('accuracy')


# we should use max_depth = 15

```

```

☐ Text(0, 0.5, 'accuracy')

```



```

[71] # prune the decision tree by using max_depth = 15

# call the pruned model
tree_model_prune = tree.DecisionTreeClassifier(max_depth = 15)

# fit the model with our data
tree_fit_prune = tree_model_prune.fit(X_train, Y_train)

# predict using our model
tree_pred_prune = tree_fit_prune.predict(X_test)


# accuracy score of pruned model
print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(tree_pred_prune, Y_test)))

Accuracy:    0.983172

```

```
# overall image of the pruned decision tree
print(tree.plot_tree(tree_model_prune))
```

```
[Text(0.5856309557174888, 0.96875, 'x[77] <= 0.5\ngini = 0.487\nsamples = 40170\nvalue = [16830, 23340]'), Text(0.396
```



From above, we can see that the decision tree clearly became less deep. This made the model less complex. However, we noticed a drop in accuracy score after pruning the tree so there is a tradeoff between computational cost and performance.

## ▼ Model 2: Naive Bayes

Naive Bayes is a probabilistic machine learning model used for classification tasks. It is based on Bayes' theorem, which is a fundamental principle in probability theory that describes how to update probabilities based on new evidence. The Naive Bayes model makes some simplifying assumptions about the data, including that the features are independent of each other and that each feature contributes equally to the probability of the class.

```
0s # Gaussian Naive Bayes

from sklearn.naive_bayes import GaussianNB

G_bayes_model = GaussianNB()
G_bayes_model.fit(X_train, Y_train)

G_bayes_pred = G_bayes_model.predict(X_test)

print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(G_bayes_pred, Y_test)))

Accuracy:    0.626904
```

```
0s [74] # Bernoulli Naive Bayes

from sklearn.naive_bayes import BernoulliNB

B_bayes_model = BernoulliNB()

B_bayes_model.fit(X_train, Y_train)

B_bayes_pred = B_bayes_model.predict(X_test)

print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(B_bayes_pred, Y_test)))

Accuracy:    0.785024
```

Since the normality assumption likely does not hold in these given features, the Gaussian Naive Bayes model did not yield a very high accuracy rate. In contrary, the Bernoulli Naive Bayes model had provided us with a higher accuracy rate. This is because the Bernoulli Naive Bayes model is most appropriate when working with binary or boolean data and in our prepared data frame, there were more features that involved binary values. Overall, since the data had both binary and continuous values, the Naive Bayes Model did not perform very well. Additionally, the overall model assumption that the predictors are independent could have been violated as well.



## ▼ 9. Experiment with 1 Ensemble Method

```
✓ [75] # Ensemble Methods: https://scikit-learn.org/stable/modules/ensemble.html
```

### ▼ Random Forest

```
✓ [76] # import necessary packages
      from sklearn.ensemble import RandomForestClassifier

      # call the random forest model
      forest_model = RandomForestClassifier(n_estimators = 5)

      # fit the model with data
      forest_model_fit = forest_model.fit(X_train, Y_train)

      # predict using our model
      forest_model_pred = forest_model_fit.predict(X_test)

      # accuracy score
      print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(forest_model_pred, Y_test)))

Accuracy:      0.999502
```

```
✓ [77] # now try with n_estimators = 10

      # call the random forest model
      forest_model = RandomForestClassifier(n_estimators = 10)

      # fit the model with data
      forest_model_fit = forest_model.fit(X_train, Y_train)

      # predict using our model
      forest_model_pred = forest_model_fit.predict(X_test)

      # accuracy score
      print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(forest_model_pred, Y_test)))

Accuracy:      0.999801
```

The accuracy has slightly gone up after tuning the number of estimators.

```
✓ [78] # now we experiment with the max depth

      # call the random forest model
      forest_model = RandomForestClassifier(n_estimators = 10, max_depth = 15)

      # fit the model with data
      forest_model_fit = forest_model.fit(X_train, Y_train)

      # predict using our model
      forest_model_pred = forest_model_fit.predict(X_test)

      # accuracy score
      print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(forest_model_pred, Y_test)))

Accuracy:      0.994324
```

The accuracy has slightly gone down after manipulating the max\_depth

## ▼ 10. Cross-Validation & Hyperparameter Tuning for All 3 Models

```
[79] # Cross-Validation: https://scikit-learn.org/stable/modules/cross\_validation.html  
      # Hyperparameter Tuning: https://scikit-learn.org/stable/modules/grid\_search.html
```

### ▼ Cross-Validation

```
[80] from sklearn import model_selection  
      from sklearn.model_selection import KFold  
  
      # define number of splits, seed, and let data be shuffled  
      kfold = model_selection.KFold(n_splits = 10, random_state = 42, shuffle = True)  
  
      # call cross_val_score with argument cv = kfold to get accuracy performance score  
      # of our model on the 10 fold validation data  
      # other arguments: model, training_data, training labels  
      results = model_selection.cross_val_score(log_model, train_prepared_df, label, cv = kfold)  
  
      # print the average score for your model on the 10 folds  
      print("Cross Validation Accuracy: %.2f%%" % (results.mean()*100.0))  
  
Cross Validation Accuracy: 85.68%
```

```
[81] # for tree model  
      results = model_selection.cross_val_score(tree_model, train_prepared_df, label, cv = kfold)  
  
      # print the average score for your model on the 10 folds  
      print("Cross Validation Accuracy: %.2f%%" % (results.mean()*100.0))  
  
Cross Validation Accuracy: 99.84%
```

```
[82] # for pruned tree model  
      results = model_selection.cross_val_score(tree_model_prune, train_prepared_df, label, cv = kfold)  
  
      # print the average score for your model on the 10 folds  
      print("Cross Validation Accuracy: %.2f%%" % (results.mean()*100.0))  
  
Cross Validation Accuracy: 98.31%
```

```
[83] # for bayes classifier  
      results = model_selection.cross_val_score(B_bayes_model, train_prepared_df, label, cv = kfold)  
  
      # print the average score for your model on the 10 folds  
      print("Cross Validation Accuracy: %.2f%%" % (results.mean()*100.0))  
  
Cross Validation Accuracy: 77.93%
```

```
[84] # for random forest  
      results = model_selection.cross_val_score(forest_model, train_prepared_df, label, cv = kfold)  
  
      # print the average score for your model on the 10 folds  
      print("Cross Validation Accuracy: %.2f%%" % (results.mean()*100.0))  
  
Cross Validation Accuracy: 99.43%
```

## Hyperparameter Tuning

```
[85] from sklearn.model_selection import GridSearchCV
```

### Decision Tree Model

```
[86] params = [{
    'criterion': ['gini', 'entropy', 'log_loss'],
    'splitter': ['best', 'random'],
    'max_depth': [5, 10, 20, 30]
}]

# define grid search cross validation
clf2 = model_selection.GridSearchCV(estimator = tree_model, param_grid = params, scoring = "accuracy", cv = kfold, verbose = True)

# call on the data
clf2_fit = clf2.fit(train_prepared_df, label)

# retrieve estimator which gave highest score on cross validation
best_model = clf2_fit.best_estimator_
print(best_model)

# get the mean score for your best model
print(clf2_fit.best_score_)
```

Fitting 10 folds for each of 24 candidates, totalling 240 fits  
DecisionTreeClassifier(max\_depth=30, splitter='random')  
0.9995220273973169

### Bayes Classifier

```
[87] params = [{
    'alpha': [0.5, 0.8, 1],
    'force_alpha': [True, False]
}]

# define grid search cross validation
clf3 = model_selection.GridSearchCV(estimator = B_bayes_model, param_grid = params, scoring = "accuracy", cv = kfold, verbose = True)

# call on the data
clf3_fit = clf3.fit(train_prepared_df, label)

# retrieve estimator which gave highest score on cross validation
best_model = clf3_fit.best_estimator_
print(best_model)

# get the mean score for your best model
print(clf3_fit.best_score_)
```

Fitting 10 folds for each of 6 candidates, totalling 60 fits  
BernoulliNB(alpha=0.5, force\_alpha=True)  
0.7793800526042316

### Random Forest

```
[94] params = [{
    'n_estimators': [10, 50, 100],
    'criterion': ["gini", "entropy", "log_loss"],
    'max_depth': [10, 20, 30],
    'min_samples_leaf': [1, 2, 3]
}]

# define grid search cross validation
clf4 = model_selection.GridSearchCV(estimator = forest_model, param_grid = params, scoring = "accuracy", cv = kfold, verbose = True)

# call on the data
clf4_fit = clf4.fit(train_prepared_df, label)

# retrieve estimator which gave highest score on cross validation
best_model = clf4_fit.best_estimator_
print(best_model)

# get the mean score for your best model
print(clf4_fit.best_score_)
```

Fitting 10 folds for each of 81 candidates, totalling 810 fits  
RandomForestClassifier(max\_depth=30, n\_estimators=50)  
0.9999203425263434



## 11. Report Final Results

```
[108] # Temporarily combining the training and testing, so that necessary OHE can be made
train_test_comb = imputed_train_dropped.append(imputed_test_dropped)

# as fit_transform outputs an np.array, we aim to convert to a df

# we aim to retrieve column names
ohe_encoder = OneHotEncoder(categories = "auto", handle_unknown = "ignore", sparse_output = False)
X_object = (train_test_comb.drop(["cap-diameter", "stem-height", "stem-width"], axis = 1))
ohe_encoder.fit(X_object)

# fit the full pipeline
train_test_comb_prepared = full_pipeline.fit_transform(train_test_comb)

# features that will be introduced through One Hot Encoding
feature_names_ohe = ohe_encoder.get_feature_names_out(OHE)

# convert the np.array to df (this df will not have column names)
train_test_comb_prepared_df = pd.DataFrame(train_test_comb_prepared)

# combine numerical features and features that were OHE
features_df = numerical_features + list(feature_names_ohe)

# rename df
train_test_comb_prepared_df.set_axis(features_df, axis = 1, inplace = True)

train_test_comb_prepared_df.head()
```

	cap-diameter	stem-height	stem-width	cap-shape_b	cap-shape_c	cap-shape_f	cap-shape_o	cap-shape_p	cap-shape_s	cap-shape_x	...	habitat_h	habitat_l	habitat_m	habitat_p	habitat_u	habitat_w	season_a	season_s	season_u	season_w
0	1.619462	3.076705	0.492293	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
1	1.873982	3.385311	0.601900	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	0.0
2	1.393432	3.328931	0.557061	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
3	1.412426	2.726555	0.381690	0.0	0.0	1.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0
4	1.501699	2.952075	0.503254	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0

5 rows x 118 columns

```
[109] # resplit back to training and testing
train_final = train_test_comb_prepared_df.iloc[:50213, :]
test_final = train_test_comb_prepared_df.iloc[50213:, :]
```

### Decision Tree

```
[116] # model with best hyperparameter

# call the model
tree_model_best = tree.DecisionTreeClassifier(max_depth=30, splitter='random')

# fit the model with our data
tree_best_fit = tree_model_best.fit(train_final, label)

# predict using our model
tree_best_pred = tree_best_fit.predict(test_final)
```

```
[117] # the accuracy of the model
print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(tree_best_pred, label_test)))

# the precision of the model
print("%-12s %f" % ('Precision:', metrics.precision_score(tree_best_pred, label_test)))

# the recall score of the model
print("%-12s %f" % ('Recall Score:', metrics.recall_score(tree_best_pred, label_test)))

# the F1 score of the model
print("%-12s %f" % ('F1 Score:', metrics.f1_score(tree_best_pred, label_test)))
```

Accuracy: 0.432019  
Precision: 0.792294  
Recall Score: 0.414434  
F1 Score: 0.544205

## ▼ Bayes Classifier

```
✓ [118] # model with best hyperparameter
```

0s

```
# call the bayes model
B_bayes_model_best = BernoulliNB(alpha=0.5, force_alpha=True)

# fit the best model
B_bayes_best_fit = B_bayes_model_best.fit(train_final, label)

# predict using our model
B_bayes_best_pred = B_bayes_best_fit.predict(test_final)
```

```
✓ [119] # the accuracy of the model
```

0s

```
print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(B_bayes_best_pred, label_test)))

# the precision of the model
print("%-12s %f" % ('Precision:', metrics.precision_score(B_bayes_best_pred, label_test)))

# the recall score of the model
print("%-12s %f" % ('Recall Score:', metrics.recall_score(B_bayes_best_pred, label_test)))

# the F1 score of the model
print("%-12s %f" % ('F1 Score:', metrics.f1_score(B_bayes_best_pred, label_test)))
```

```
Accuracy:      0.498066
Precision:     1.000000
Recall Score:  0.460228
F1 Score:      0.630351
```

## ▼ Random Forest Classification

```
✓ [120] # model with best hyperparameter
```

3s

```
# call the random forest model
forest_model_best = RandomForestClassifier(max_depth=30, n_estimators=50)

# fit the model with data
forest_model_best_fit = forest_model_best.fit(train_final, label)

# predict using our model
forest_model_best_pred = forest_model_best_fit.predict(test_final)
```

```
[121] # the accuracy of the model
```

```
print("%-12s %f" % ('Accuracy:', metrics.accuracy_score(forest_model_best_pred, label_test)))

# the precision of the model
print("%-12s %f" % ('Precision:', metrics.precision_score(forest_model_best_pred, label_test)))
~
# the recall score of the model
print("%-12s %f" % ('Recall Score:', metrics.recall_score(forest_model_best_pred, label_test)))

# the F1 score of the model
print("%-12s %f" % ('F1 Score:', metrics.f1_score(forest_model_best_pred, label_test)))
```

```
Accuracy:      0.459285
Precision:     0.956522
Recall Score:  0.439478
F1 Score:      0.602250
```